

Database Administration:

The Complete Guide to Practices and Procedures

Chapter 3 Database Design



Agenda

- From Logical Model to Physical Database
- Database Performance Design
- Denormalization
- Views
- Data Definition Language
- Temporal Data Support
- Questions

Physical Database Design Requirements

- In-depth knowledge of the database objects supported by the DBMS and the physical structures and files required to support those objects
- Details regarding the manner in which the DBMS supports indexing, referential integrity, constraints, data types, and other features that augment the functionality of database objects
- Detailed knowledge of new and obsolete features for particular versions or releases of the DBMS
- Knowledge of the DBMS configuration parameters that are in place
- Data definition language (DDL) skills to translate the physical design into actual database objects

Basic Physical ROTs

- Avoid using default settings
 - They are rarely the best setting
 - It is better to know and explicitly state the actual setting you desire in each case
- Synchronize the logical and physical models
 - Always map changes in one to the other
- Performance before aesthetics
- Meaning: prefer fully normalized
- but deviate when necessary to
 - achieve performance goals
- Almost never say *always* or *never*




Transforming Logical to Physical

- Translation of Logical Model to Physical Database
 - Create DDL
 - Entities to Tables, Attributes to Columns, Relationships and Keys to DB2 RI and Indexes, etc.
 - ...but differences **CAN** and **WILL** occur
- Create Storage Structures for Database
 - Files for data and indexes
 - Partitioning
 - Clustering
 - Placement
 - Order of columns



Transform Entities to Tables

- First general step:
 - Map each entity in the logical data model to a table in the database
- Things may, or may not be that easy
 - Denormalization?



Transform Attributes to Columns

- Attributes become columns
- Transform Domains to Data Types
 - Commercial DBMSs do not support domains
 - Date Type and Length
 - Variable or Fixed Length
 - Choose wisely; impacts data quality
 - Constraints
 - Null

Data Types

- CHAR / VARCHAR
- CLOB
- DBCLOB
- BLOB
- GRAPHIC / VARGRAPHIC
- DATE
- TIME
- DATETIME / TIMESTAMP
- XML
- BIGINT
- INTEGER
- SMALLINT
- MONEY
- BINARY
- DECIMAL
- FLOAT
 - REAL
 - DOUBLE

Nulls

INVEN_LOC_TAB

WAREHSE_NO SMALLINT	BIN_NO SMALLINT	PROD_NO CHAR(5)	PROD_QTY INT
1	1	A100	???
1	2	A150	2000
1	3	B167	30
2	1	A100	775
2	2	D400	1585

NULL:

Has no value

Is not = anything

Is not < anything

Is not > anything

Is not = NULL

Is an UNKNOWN value

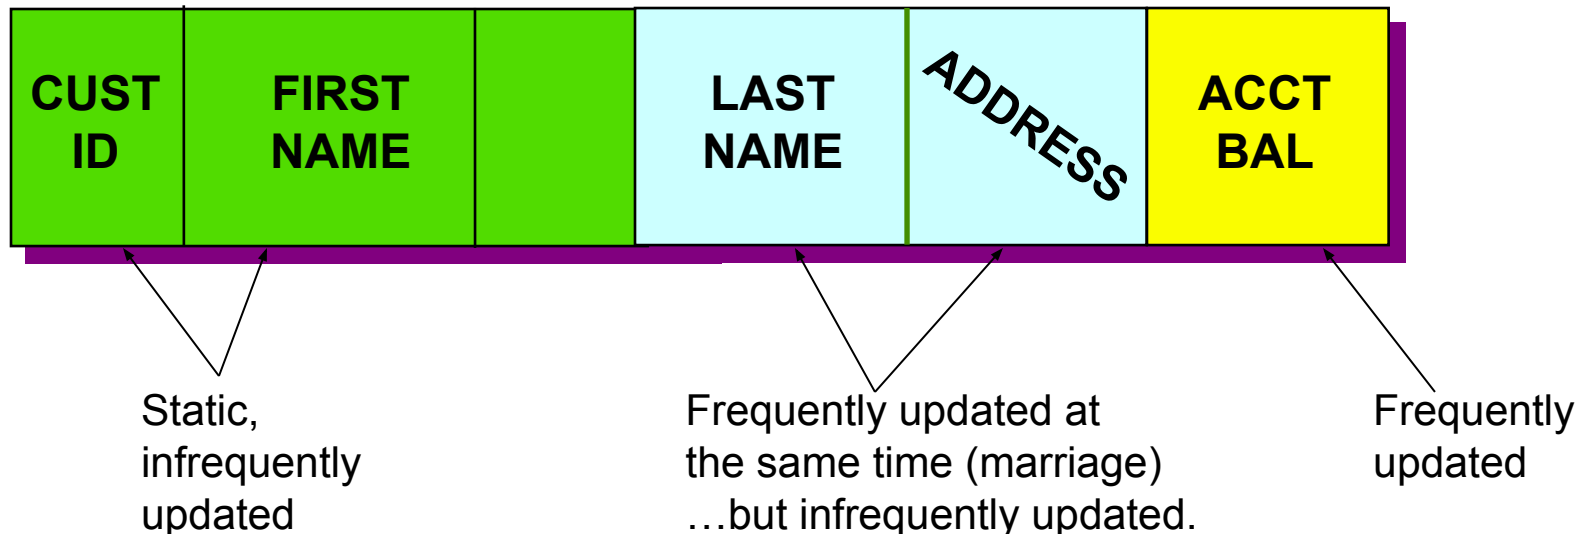
ATTRIBUTE QUALIFIER	DESCRIPTION
DEFAULT NULL	When no value is provided, DB2 automatically assigns nulls to the Column
NOT NULL	The column must always contain a <u>value</u> , whether a default or explicitly provided

DEFAULT

ATTRIBUTE QUALIFIER	DESCRIPTION	DEFAULT USED
(WITH) DEFAULT	When no value is provided, DB2 automatically assigns an appropriate default	Numeric Data: ZEROS
		Character Data: SPACES
		Variable Data: ZERO LENGTH
		Chronological Data: CURRENT DATE CURRENT TIME CURRENT TIMESTAMP
(WITH) DEFAULT value		Constant USER CURRENT SQLID NULL

Column Ordering

- Sequence columns based on logging. For example:
 - Infrequently updated non-variable columns first
 - Static (infrequently updated) variable columns
 - Frequently updated columns last
 - Frequently modified together, place next to each other



Relationships and Keys

- Use the primary key as assigned in the logical modeling phase for the physical PK
- Other considerations:
 - Length of the key
 - Surrogate key
 - ROWID / SEQUENCE / Identity
- Build referential constraints for all relationships
 - Foreign keys

Build Physical Structures

- Table Spaces
- DBSpaces
- Data Spaces
- Filegroups

Determine Row Size

TABLE Name: _____

Column Name	Data Type	Length (in bytes)	Variable (add 2)	Null (add 1)	TOTAL

Data Lengths:	Data Length:	
	Row Overhead:	+8
	Physical Row Length:	


INTEGER	4	FLOAT (DBL PRECISION)	8
DECIMAL PACKED FORMAT		DATE	4
SMALLINT	2	TIME	3
FLOAT (REAL)	4	TIMESTAMP	10

Storage Planning

- Start by determining how many rows are required
- Calculate the row size (*discussed earlier*)
- Figure out the number of rows per block/page
- Multiple #rows/page by the page size
- This gives you the size of the object
- Except for free space...


Freespace Alternatives

- ♦ Freespace is considered only at LOAD or REORG time
- ♦ The freespace options:
PCTFREE
FREEPAGE
- ♦ For Example:



Type of Files

- Data / Index
 - Both require storage
- Raw Files
 - Can be used to bypass the O/S
- Solid State Devices
 - For performance-critical objects
- *Compression*



Database Performance Design

- Designing Indexes
 - Partitioning
 - Clustering
- Hashing
- Interleaving Data

Designing Indexes

Index Advantages

Optimize data access:

- DBMS decides whether or not to use an index
- DBMS maintains all indexes (*modifications incur cost*)
- Table scans can be avoided through index usage
- Recommended on foreign key columns to speed RI access
- Indexes can minimize sorting
- There can be multiple indexes per table to suit the way data is processed
- Create indexes based on workload (not tables)
- If all columns are in the index you can get index-only access (IXO)


Guarantee uniqueness:

- Can be used to ensure uniqueness of column values
- Required on primary key column as part of referential integrity implementation

Implement clustering:

- Indexes can be used for clustering; that is, maintaining the rows physically on disk in the sequence of the column values in the index

B-Tree Index



...to the data in the table.

Bitmap Index

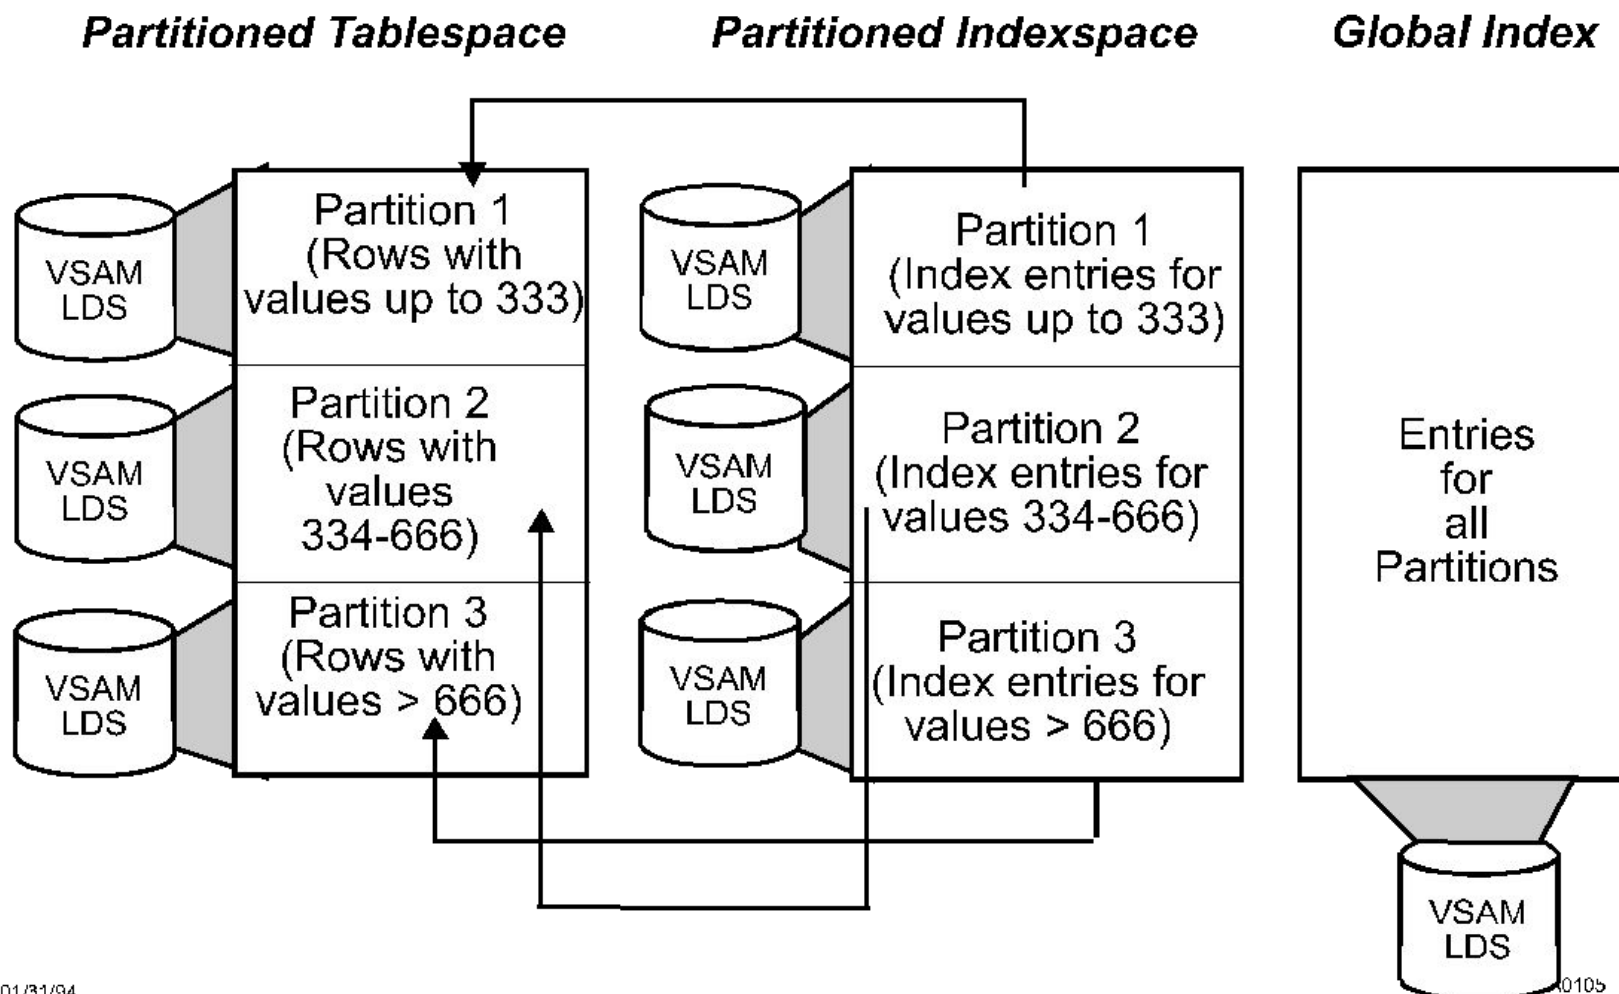
Identifier	Gender	Bitmap
1	Female	0110000010
2	Male	1000011101
3	Unknown	0001100000

Other Types of Indexes

- Reverse Key Index
 - a b-tree index where the order of bytes of each indexed column is reversed; helps with hot spots
- Partitioned Index
 - a b-tree index specifying how to break up the index (and perhaps the underlying table) into separate chunks, or partitions; to enhance performance and availability
- Ordered Index

Partitioning

Partitioned Tablespace



Clustering

Clustering Index

NON-UNIQUE
CLUSTERING INDEX
ORDR_DATE

SALES_ORDER_TAB

SALES_ORDER_NO INT	ORDR_DATE DATE	CUST_NO SMALL INT	SALES_HIST_CUST_NO SMALL INT	ORDR_AMT DEC(9,2)
1	1997-09-15		2	1923.45
3	1997-09-23		1	2407.53
4	1997-09-23		10	57613.89
5	1997-09-29	3	5	67000.00
6	1997-10-04	2		42345.88
7				122345.61
8				23007.34
9				9823.55
10				223019.27
11				78780.99

CLUSTERING...

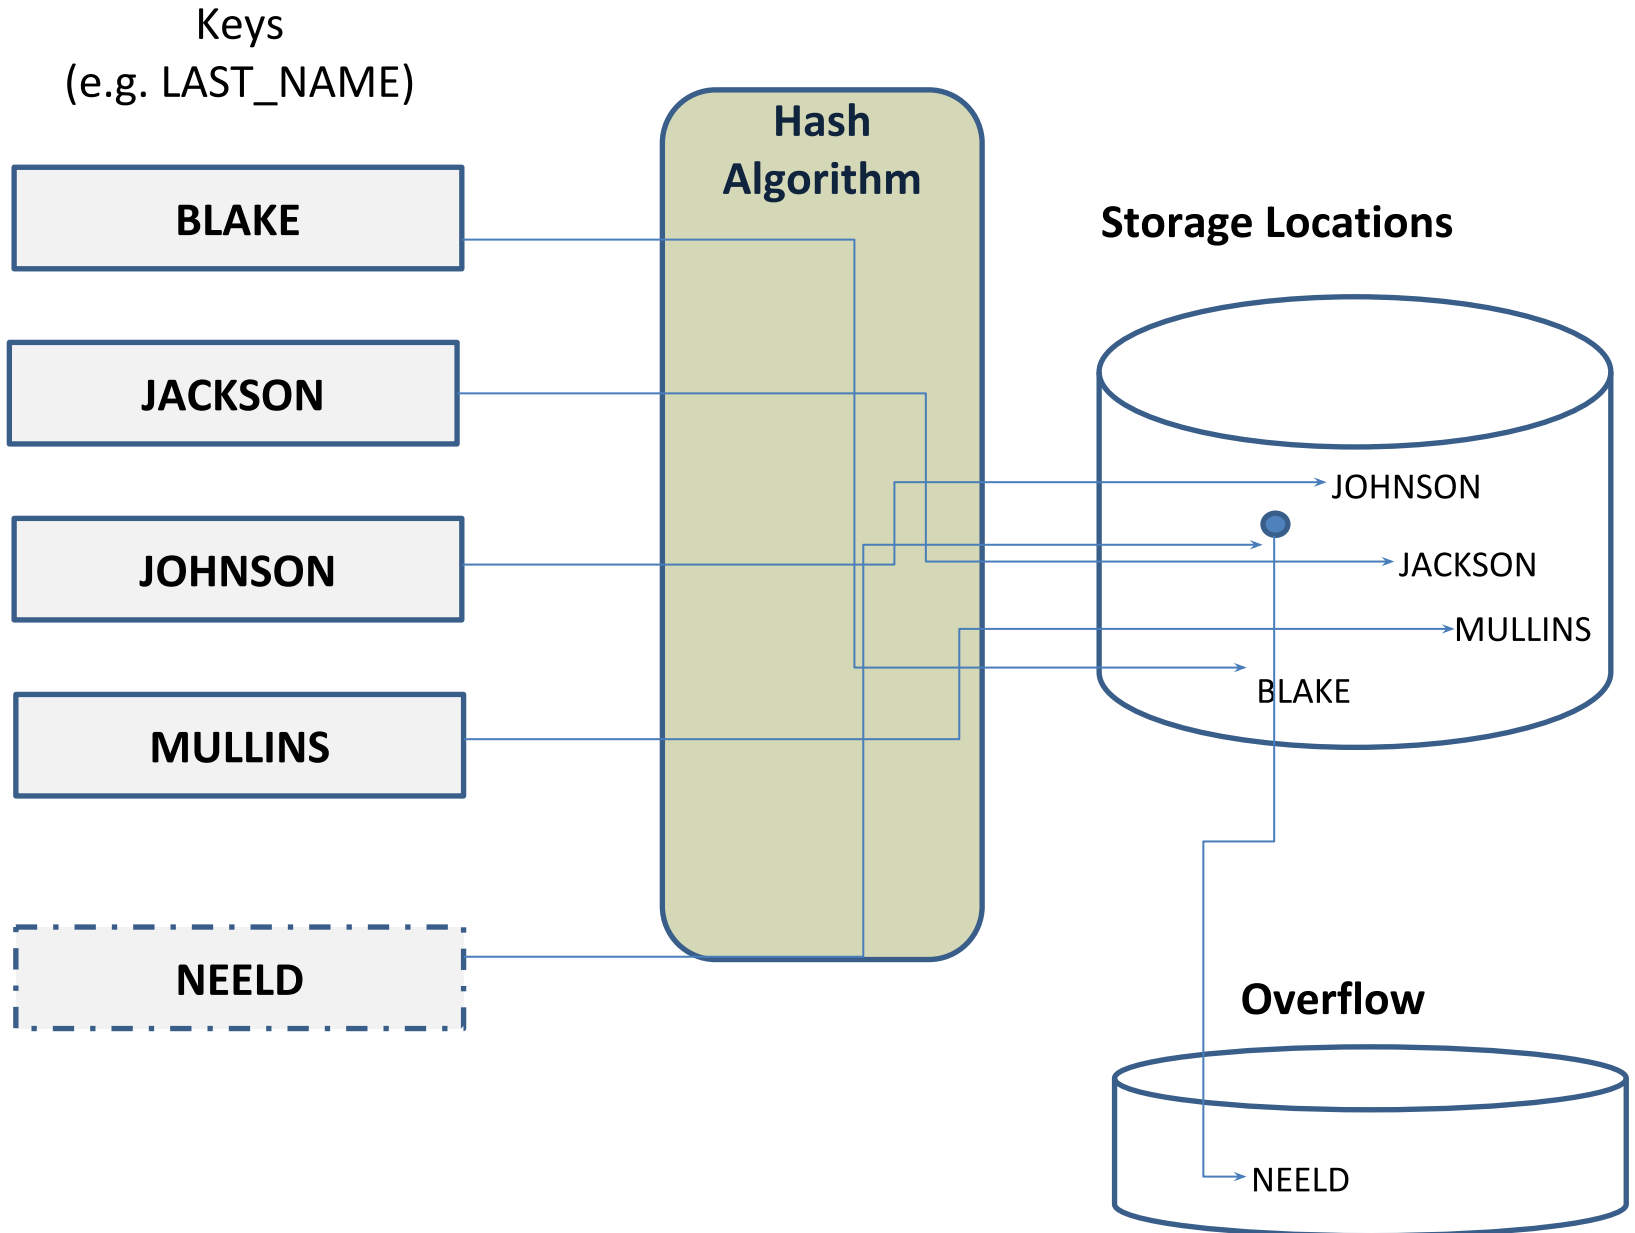
GIVES I/O

MORE


"BANG FOR THE BUCK!!"

ORDR_NO	DATE
000000125	19970923
000004946	19970923
000000013	19971002
000000615	19971002
000008885	19971002
000074155	19971004

Hashing



Interleaving Data



Legend

○ Table 1

● Table 2

Denormalization

- **Prejoined Tables** - when the cost of joining is prohibitive
- **Report Tables** - for specialized critical reports (e.g. CEO)
- **Mirror Tables** - when two types of environments require concurrent access to the same data (**OLTP vs DSS**)
- **Split Tables** - when distinct groups/apps use different parts of the same table
 - Splitting **columns** across two tables for long variable character columns.
- **Combined Tables** - to eliminate one-to-one relationships
- **Redundant Data** - to reduce the number of joins for a single column (e.g. definitional, CA to California)
- **Repeating Groups** - to reduce overall I/O (& possibly DASD)
- **Derivable Data** - to eliminate calculations & aggregations
- **Speed Tables** - to support hierarchies
- **Physical Implementation Needs** – e.g.) to reduce page size

When to Denormalize

The only reason to denormalize, ever:

- To achieve optimal **performance**!
- If the database design achieve satisfactory performance fully normalized, then there is no need to denormalize.

You should always consider the following issues before denormalizing.

- Can the system achieve acceptable performance *without* denormalizing?
- Will the performance of the system *after* denormalizing still be unacceptable?
- Will the system be less reliable due to denormalization?

Denormalization Administration

The decision to denormalize should never be made lightly, because it can cause integrity problems and involve a lot of administration overheads.

Additional administration tasks include:

- Documenting every denormalization decision
- Ensuring that all data remains valid and accurate
- Scheduling data migration and propagation jobs
- Keeping end users informed about the state of the tables
- Analyzing the database periodically to decide whether denormalization is still required

Normalized vs. Denormalized

Normalized Tables:

The Goal!




- ♦ More Tables
- ♦ Fewer Columns per Table
- ♦ Fewer Rows per Table
- ♦ Less Redundancy
- ♦ More Joins
- ♦ Update Efficient

Denormalized Tables:

- ♦ Fewer Tables
- ♦ More Columns per Table
- ♦ More Rows per Table
- ♦ Greater Redundancy
- ♦ Fewer Joins
- ♦ Read Efficient

Views



What is a View?

Course Table = TCRSE An actual table:

CRSE_NUM	CRSE_CLASS	CRSE_INST	CRSE_NAME	CRSE_DATE	CRSE_TIME	CRSE_DAY	CRSE_TYPE
INT	CHAR(6)	INT	CHAR(20)	DATE	TIME	CHAR(2)	CHAR(1)
26	DB2004	200	SYSTEMS ADMIN	1990-11-24	10.00.00	TU	A
27	DB2001	--	INTRO TO DB2	1990-11-30	09.00.00	MO	A
28	DB2002	300	SQL PROGRAMMING	1990-11-30	09.00.00	MO	B
29	PHY001	500	RELATIVITY	1990-10-02	09.00.00	WD	A
30	DB2003	--	DATABASE DESIGN	1990-12-07	09.00.00	MO	B
31	DB2001	800	INTRO TO DB2	1990-12-07	09.00.00	MO	A
32	LIT001	600	ENTREPRENEUR	1990-12-09	10.00.00	WD	A
33	DB2001	100	INTRO TO DB2	1990-12-09	09.00.00	WD	B

Course View = VCRSE

CRSE_NUM	CRSE_NAME	CRSE_DATE	CRSE_TIME	CRSE_DAY	CRSE_TYPE
INT	CHAR(20)	DATE	TIME	CHAR(2)	CHAR(1)
26	SYSTEMS ADMIN	1990-11-24	10.00.00	TU	A
27	INTRO TO DB2	1990-11-30	09.00.00	MO	A
28	SQL PROGRAMMING	1990-11-30	09.00.00	MO	B
29	RELATIVITY	1990-10-02	09.00.00	WD	A
30	DATABASE DESIGN	1990-12-07	09.00.00	MO	B
31	INTRO TO DB2	1990-12-07	09.00.00	MO	A
32	ENTREPRENEUR	1990-12-09	10.00.00	WD	A
	INTRO TO DB2	1990-12-09	09.00.00	WD	B

A logical
view of that
table

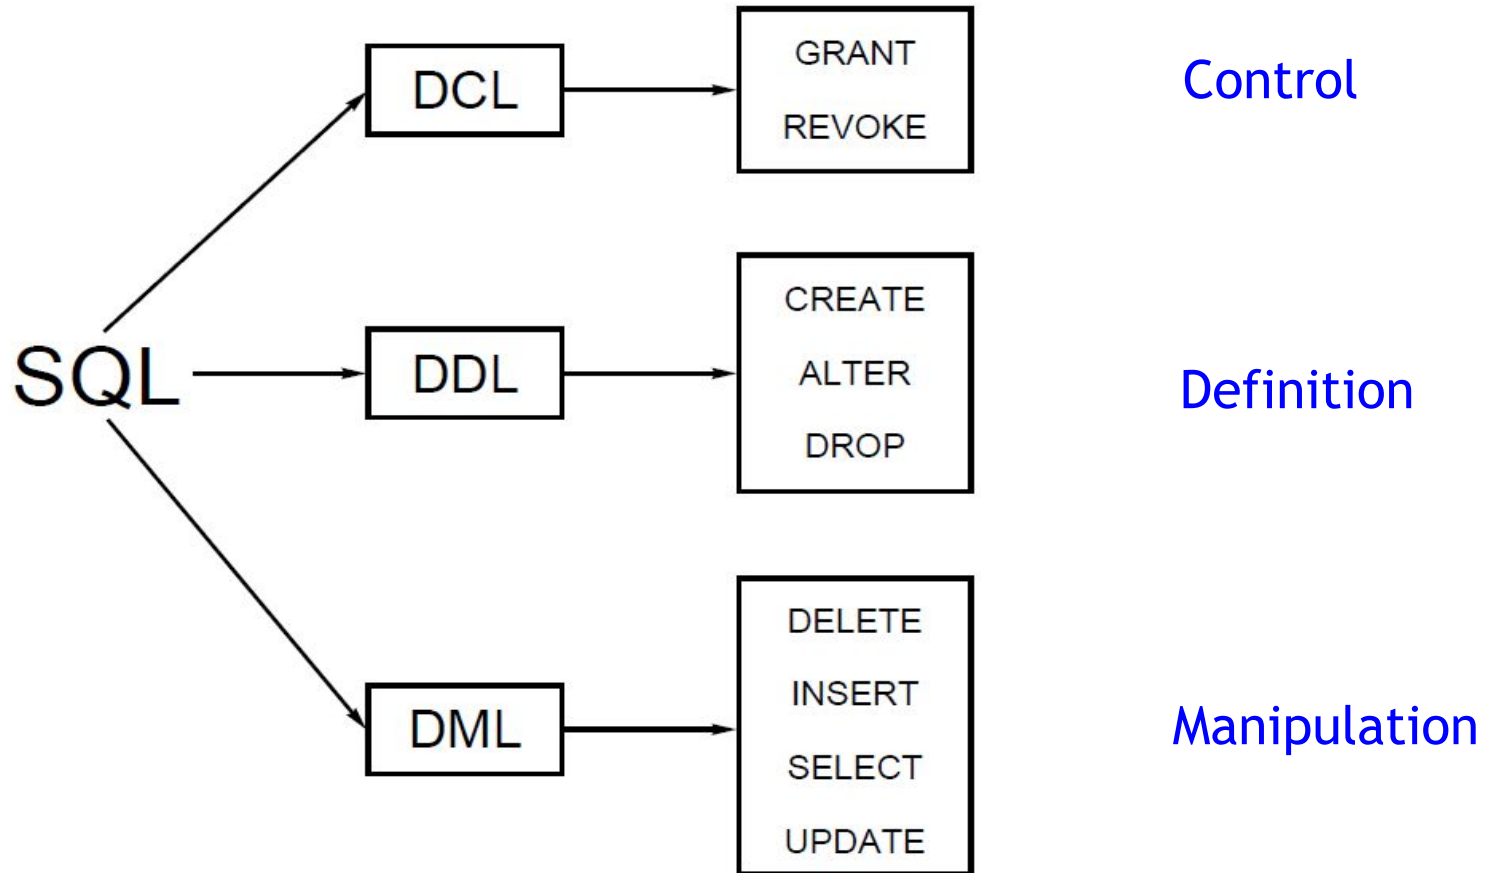
**Course list
'View'**

View Usage Rules

- Security - row and column level
- Access - efficient access paths
- Data Derivation - put the calculations in the view
- Mask Complexity - hide complex SQL from users
- Rename a Table
- Column Renaming - table with better column names (easier to use than AS)
- Synchronize all views with base tables...

DO NOT USE ONE VIEW PER BASE TABLE!

Types of SQL



Questions

